

Deep Learning Homework 2

1. Record the training and testing accuracy, plot the training loss curve and training accuracy curve in the report.

Softmax for MNIST Classification

Training

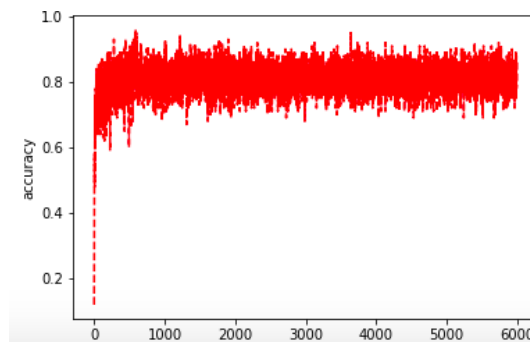
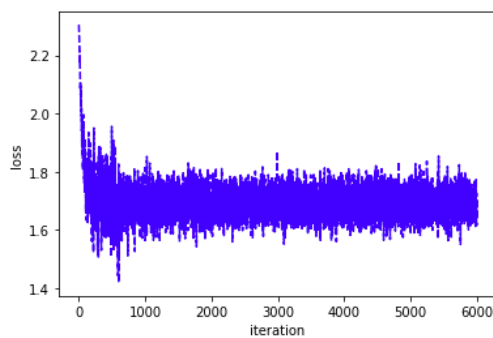
Epoch [9][10]	Batch [0][600]	Training Loss 1.6400	Accuracy 0.9200
Epoch [9][10]	Batch [100][600]	Training Loss 1.6175	Accuracy 0.7900
Epoch [9][10]	Batch [200][600]	Training Loss 1.6477	Accuracy 0.8600
Epoch [9][10]	Batch [300][600]	Training Loss 1.6287	Accuracy 0.8200
Epoch [9][10]	Batch [400][600]	Training Loss 1.6912	Accuracy 0.8000
Epoch [9][10]	Batch [500][600]	Training Loss 1.6639	Accuracy 0.8200

Testing

```
accuracy = correct * 1.0 / test_size  
print('Test Accuracy: ', accuracy)
```

Test Accuracy: 0.8237

Plot



1. MLP for MNIST Classification

1.1 MLP with Euclidean Loss and Sigmoid Activation Function

Training

Epoch [19]	Average training loss 0.0925	Average training accuracy 0.9244
Epoch [19]	Average validation loss 0.0784	Average validation accuracy 0.9454

Testing

```
test(sigmoidMLP, criterion, data_test, batch_size, disp_freq)
```

Testing...
The test accuracy is 0.9296.

1.2 MLP with Euclidean Loss and ReLU Activation Function

Training

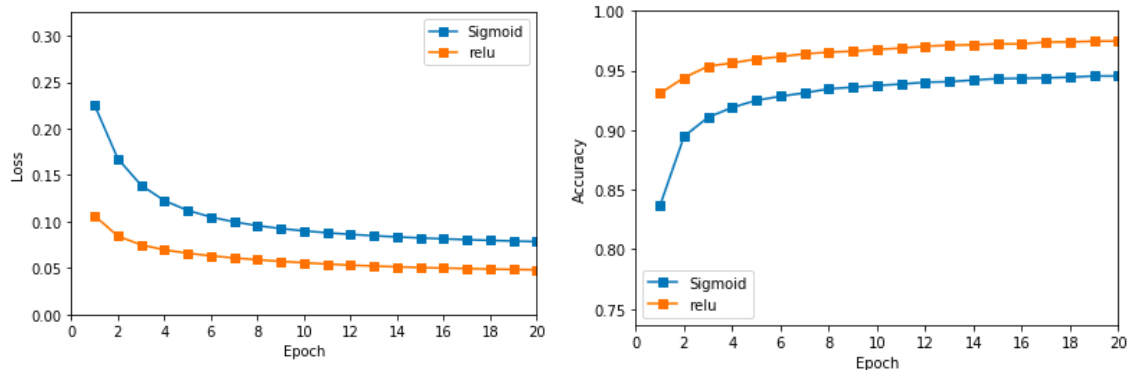
Epoch [19]	Average training loss 0.0509	Average training accuracy 0.9674
Epoch [19]	Average validation loss 0.0480	Average validation accuracy 0.9748

Testing

```
test(reluMLP, criterion, data_test, batch_size, disp_freq)
```

Testing...
The test accuracy is 0.9656.

Plot



2. MLP with Softmax Cross-Entropy Loss

2.1 MLP with Softmax Cross-Entropy Loss and Sigmoid Activation Function

Training

Epoch [19]	Average training loss 1.8726	Average training accuracy 0.6926
Epoch [19]	Average validation loss 1.8559	Average validation accuracy 0.7236

Testing

```
test(sigmoidMLP, criterion, data_test, batch_size, disp_freq)
```

Testing...
The test accuracy is 0.7050.

2.2 MLP with Softmax Cross-Entropy Loss and ReLU Activation Function

Training

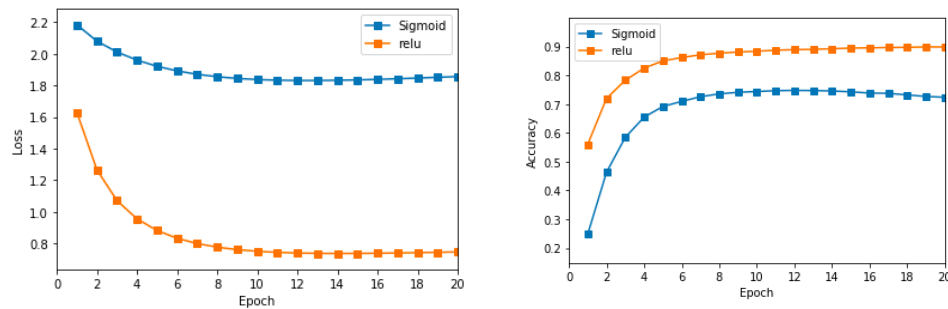
Epoch [19]	Average training loss 0.8175	Average training accuracy 0.8584
Epoch [19]	Average validation loss 0.7464	Average validation accuracy 0.8984

Testing

```
test(reluMLP, criterion, data_test, batch_size, disp_freq)
```

Testing...
The test accuracy is 0.8710.

Plot



Two-hidden-layer MLP

Softmax Cross-Entropy Loss and ReLU/ReLU Activation Function

Training

Epoch [19]	Average training loss 0.8431	Average training accuracy 0.8558
Epoch [19]	Average validation loss 0.7814	Average validation accuracy 0.8880

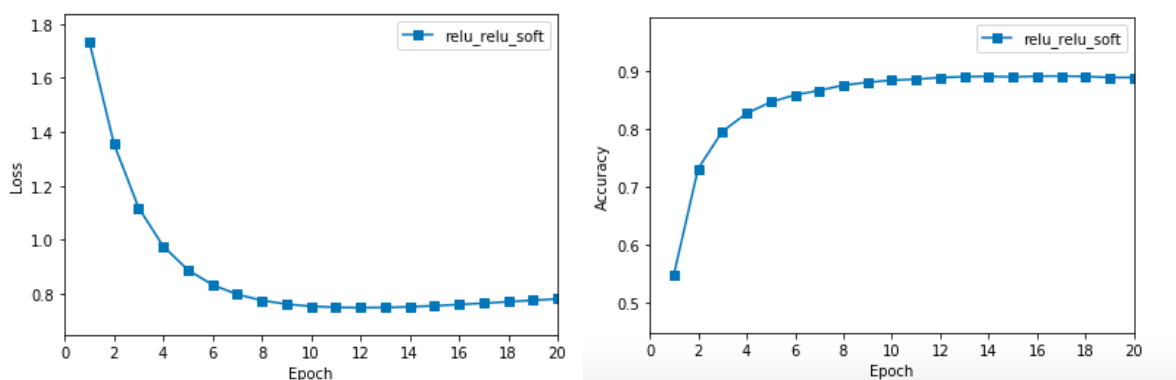
Testing

```
test(relu_relu_soft_MLP, criterion, data_test, batch_size, disp_freq)
```

Testing...

The test accuracy is 0.8643.

Plot

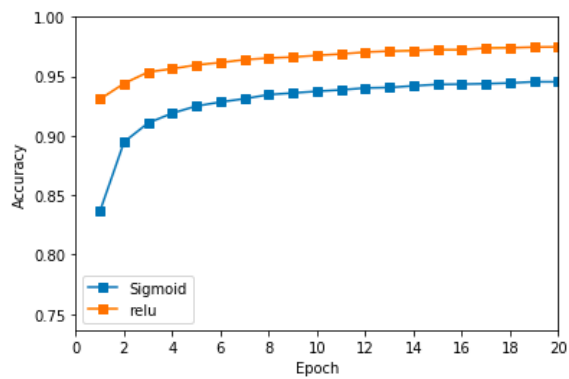
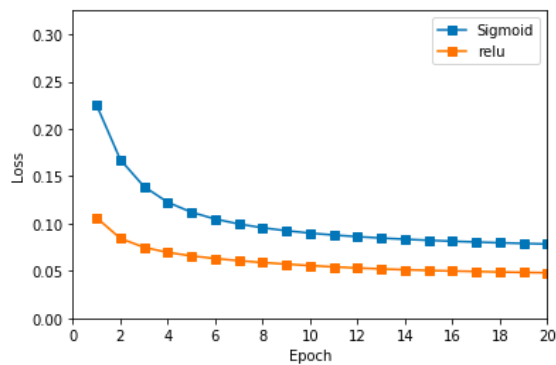


2. The given hyperparameters maybe performed not very well. You can modify the hyperparameters by your own, and observe how does these hyperparameters affect the classification performance. Write down your observation and record these new results in the report.

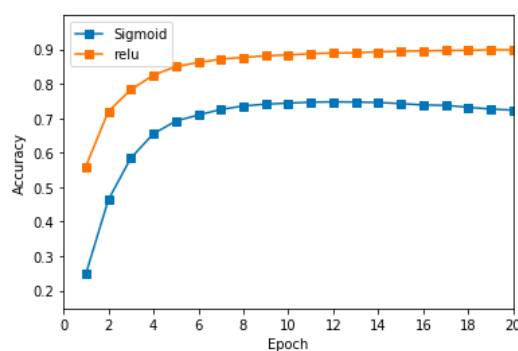
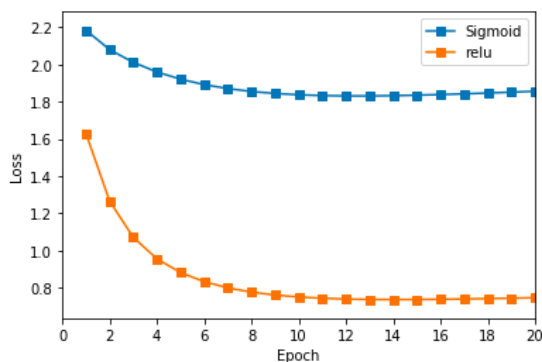
My hyperparameters performed reasonably well so I didn't have the need to change it.

3. Compare the difference of results when using Sigmoid and ReLU as activation function (you can discuss the difference from the aspects of training time, convergence and accuracy).

Euclidean Loss



Softmax Cross-Entropy Loss



3. As you can see from above ReLU performed better in both Euclidean and Softmax Cross Entropy Loss than Sigmoid.

4. Compare the difference of results when using Euclidean Loss and Softmax Cross-Entropy Loss as loss function

4. The highest accuracy rate 0.9656 was achieved by using MLP with Euclidean Loss and ReLU Activation function. The lowest was 0.7050 with Softmax Cross-Entropy Loss and Sigmoid Activation Function. When we compare the results between Euclidean and Softmax Cross Entropy we can see that Euclidean performed better with both Sigmoid and ReLU.

5. Construct a MLP with two hidden layers (choose the number of hidden units by your own), using any activation function and loss function. Also, compare the difference of results between one-layer structure and two layers structure.

I chose to compute the MLP two hidden layers with Softmax Cross-Entropy Loss and ReLU/ReLU Activation Function. As a result, it performed better than our worst result but at the end it wasn't the best but overall it performed really well and maybe if I had played with the hyperparameters or tried a different two hidden layer it would have performed better.