

1 Clustering: Mixture of Multinomials(2 points)

1.1 MLE for multinomial(1 point)

Derive the maximum-likelihood estimator for the parameter $\mu = (\mu_i)_{i=1}^d$ of a multinomial distribution:

$$P(x|\mu) = \frac{n!}{\prod_i x_i!} \prod_i \mu_i^{x_i}, \quad i = 1, \dots, d \quad (1)$$

where $x_i \in \mathbb{N}$, $\sum_i x_i = n$ and $0 < \mu_i < 1$, $\sum_i \mu_i = 1$.

Date: . . .

1.1 MLE

$$= \frac{n!}{\prod_i x_i!} \prod_i \mu_i^{x_i}$$

$$= \log \left(\frac{n!}{\prod_i x_i!} \prod_i \mu_i^{x_i} \right)$$

$$= \log(n!) - \log(\prod_i x_i!) + \log(\prod_i \mu_i^{x_i})$$

$$= \log(n!) - \sum_i \log(x_i!) + \sum_i x_i \log(\mu_i)$$

$$s.t \rightarrow x_i \in \mathbb{N}, \sum_i x_i = n$$

$$= \log(n!) - \sum_i \log(x_i!) + n \log(\mu_i)$$

Lagrange Method S.T $\rightarrow \sum_i u_i = 1 \rightarrow \sum_i -1 = 0$

$$L(u_i, x_i, \lambda) = \log(n!) - \sum_i \log(x_i!) + n \log(u_i) - \lambda (\sum u_i - 1)$$

$$\frac{dL}{du_i} = 0$$

$$= \frac{d \log(n!)}{du_i} - \frac{d(\sum_i \log(x_i!))}{du_i} + \frac{d(n \log(u_i))}{du_i} - d(\lambda (\sum u_i - 1))$$

$$= 0 + 0 + \frac{n}{u_i} - \lambda \cdot 1 = 0$$

$$= u_i = \frac{n}{\lambda}$$

$$\text{S.T} \rightarrow \sum_i u_i = 1$$

$$\frac{n}{\lambda} = 1 \Leftrightarrow n = \lambda$$

$$= u_i = \frac{x_i}{n}$$

$$\text{Answer} = u_i = \frac{x_i}{n}$$

1.2 EM for mixture of multinomials(1 point)

No.

Date.

EM AlgorithmE-StepFor every document(d), use $\lambda_d = p(c_d | \pi^{(t)}, u_{ik}^{(t)})$

$$p(c_d = k | \pi^{(t)}, u_{ik}^{(t)}) = \pi_k \prod_{w=1}^W u_{k,w}^{n_w(x_d)}$$

M-Step

Cluster label distribution:

$$u_{k,w}^{(t+1)} = \frac{\sum_{d=1}^D \lambda_{d,k} n_w(x_d)}{\sum_{d=1}^D \lambda_{d,k} n(x_d)}$$

Cluster Proposition:

$$\pi_k^{(t+1)} = \frac{\sum_{d=1}^D \lambda_{d,k}}{D}$$

2 PCA(2 points)

2.1 Minimum Error Formulation(2 points)

No.

Date.

2.1 PCA

$$J = \frac{1}{N} \sum_{n=1}^N \left\| x_n - \sum_{i=1}^d z_{ni} u_i + \sum_{i=d+1}^p b_i u_i \right\|^2$$

$$\frac{dJ}{dz_{ni}} = 0 \Leftrightarrow \frac{1}{N} (x_n^T u_i - z_{ni}) = 0$$

Result

$$\boxed{\rightarrow z_{ni} = x_n^T u_i}$$

$$\frac{dJ}{db_i} = 0 \Leftrightarrow \frac{1}{N} \sum_{n=1}^N (x_n^T u_i - b_i) = 0 \Leftrightarrow \bar{x}^T u_i - b_i = 0$$

Result

$$\boxed{\rightarrow b_i = \bar{x}^T u_i}$$

3 Reinforcement Learning(1 point)

3.1 Value Iteration(1 point)

R(s,a)	non	light	heavy		Q Table	non	light	heavy
non-send	\$ 0	\$ 20	\$ 100		non-send	-1	10	50
send	\$ (-1)	\$ 15	\$ 80		send	-2	35	40
					(After 1 Iteration)			
P(s' s, a_no-send)	non	light	heavy		Q Table	non	light	heavy
non	0.95	0.05	0		non-send	0.72	45.695	137.71
light	0.2	0.75	0.05		send	27.935	52.72	120.68
heavy	0.1	0.2	0.7					
P(s' s, a_send)	non	light	heavy		Calculation = 0 + 0.9 x (0.95 x (-1) + 0.05 x (35) + 0 x (50))			
non	0.1	0.85	0.05					
light	0.1	0.2	0.7					
heavy	0.05	0.15	0.8					

R(s,a)	non	light	heavy		Q Table	non	light	heavy
non-send	\$ 0	\$ 20	\$ 100		non-send	-1	10	50
send	\$ (-1)	\$ 15	\$ 80		send	-2	35	40
					(After 1 Iteration)			
P(s' s, a_no-send)	non	light	heavy		Q Table	non	light	heavy
non	0.95	0.05	0		non-send	0.72	45.695	137.71
light	0.2	0.75	0.05		send	27.935	52.72	120.68
heavy	0.1	0.2	0.7					
P(s' s, a_send)	non	light	heavy		Calculation = 20 + 0.9 x (0.2 x (-1) + 0.75 x (35) + 0.05 x (50))			
non	0.1	0.85	0.05					
light	0.1	0.2	0.7					
heavy	0.05	0.15	0.8					

R(s,a)	non	light	heavy		Q Table	non	light	heavy
non-send	\$ 0	\$ 20	\$ 100		non-send	-1	10	50
send	\$ (-1)	\$ 15	\$ 80		send	-2	35	40
					(After 1 Iteration)			
P(s' s, a_no-send)	non	light	heavy		Q Table	non	light	heavy
non	0.95	0.05	0		non-send	0.72	45.695	137.71
light	0.2	0.75	0.05		send	27.935	52.72	120.68
heavy	0.1	0.2	0.7					
					Calculation = 15 + 0.9 x (0.1 x (-1) + 0.2 x (35) + 0.7 x (50))			
P(s' s, a_send)	non	light	heavy					
non	0.1	0.85	0.05					
light	0.1	0.2	0.7					
heavy	0.05	0.15	0.8					

R(s,a)	non	light	heavy		Q Table	non	light	heavy
non-send	\$ 0	\$ 20	\$ 100		non-send	-1	10	50
send	\$ (-1)	\$ 15	\$ 80		send	-2	35	40
(After 1 Iteration)								
P(s' s, a_no-send)	non	light	heavy		Q Table	non	light	heavy
non	0.95	0.05	0		non-send	0.72	45.695	137.71
light	0.2	0.75	0.05		send	27.935	52.72	120.68
heavy	0.1	0.2	0.7		Calculation = $100 + 0.9 \times (0.1 \times (-1) + 0.2 \times (35) + 0.7 \times (50))$			
P(s' s, a_send)	non	light	heavy					
non	0.1	0.85	0.05					
light	0.1	0.2	0.7					
heavy	0.05	0.15	0.8					

R(s,a)	non	light	heavy		Q Table	non	light	heavy
non-send	\$ 0	\$ 20	\$ 100		non-send	-1	10	50
send	\$ (-1)	\$ 15	\$ 80		send	-2	35	40
(After 1 Iteration)								
P(s' s, a_no-send)	non	light	heavy		Q Table	non	light	heavy
non	0.95	0.05	0		non-send	0.72	45.695	137.71
light	0.2	0.75	0.05		send	27.935	52.72	120.68
heavy	0.1	0.2	0.7		Calculation = $80 + 0.9 \times (0.05 \times (-1) + 0.15 \times (35) + 0.8 \times (50))$			
P(s' s, a_send)	non	light	heavy					
non	0.1	0.85	0.05					
light	0.1	0.2	0.7					
heavy	0.05	0.15	0.8					

4 Deep Generative Models: Class-conditioned VAE(5 points)

Findings

I have found that after comparing 10 and 20 epochs. 20 was clearer on the edges and more understandable.

Here are the comparison photos:

10



20



output of test lower bound and Log Likelihood:

```
Epoch 1 (14.2s): Lower bound = -173.15237426757812
Epoch 2 (13.1s): Lower bound = -125.95645904541016
Epoch 3 (13.1s): Lower bound = -115.4300537109375
Epoch 4 (13.3s): Lower bound = -111.14623260498047
Epoch 5 (17.0s): Lower bound = -108.410400390625
Epoch 6 (18.3s): Lower bound = -106.51114654541016
Epoch 7 (18.1s): Lower bound = -105.159423828125
Epoch 8 (13.0s): Lower bound = -104.07119750976562
Epoch 9 (13.5s): Lower bound = -103.14167022705078
Epoch 10 (13.5s): Lower bound = -102.26787567138672
>>> TEST (379.8s)
>> Test lower bound = -101.43687438964844
>> Test log likelihood (IS) = -96.09476470947266
Epoch 11 (14.1s): Lower bound = -101.62956237792969
Epoch 12 (13.1s): Lower bound = -100.99659729003906
Epoch 13 (12.9s): Lower bound = -100.55213165283203
Epoch 14 (13.1s): Lower bound = -100.00975036621094
Epoch 15 (12.9s): Lower bound = -99.74927520751953
Epoch 16 (12.9s): Lower bound = -99.3624267578125
Epoch 17 (12.8s): Lower bound = -99.09620666503906
Epoch 18 (12.9s): Lower bound = -98.78494262695312
Epoch 19 (13.0s): Lower bound = -98.54756164550781
Epoch 20 (13.2s): Lower bound = -98.28917694091797
>>> TEST (402.6s)
>> Test lower bound = -98.56088256835938
>> Test log likelihood (IS) = -93.06107330322266
```

Appendix Code

```
#!/usr/bin/env python
# -*- coding: utf-8 -*-
```

```
from __future__ import absolute_import
from __future__ import print_function
from __future__ import division
import os
import time
```

```
import tensorflow as tf
from six.moves import range
import numpy as np
import zhuan as zs
```

```
from examples import conf
from examples.utils import dataset, save_image_collections
```

```
@zs.meta_bayesian_net(scope="gen", reuse_variables=True)
def build_gen(x_dim, z_dim, n, n_particles=1):
    bn = zs.BayesianNet()
    z_mean = tf.zeros([n, z_dim])
    z = bn.normal("z", z_mean, std=1., group_ndims=1, n_samples=n_particles)
    h = tf.layers.dense(z, 500, activation=tf.nn.relu)
    h = tf.layers.dense(h, 500, activation=tf.nn.relu)
    x_logits = tf.layers.dense(h, x_dim)
    bn.deterministic("x_mean", tf.sigmoid(x_logits))
    bn.bernoulli("x", x_logits, group_ndims=1)
    return bn
```

```
@zs.reuse_variables(scope="q_net")
def build_q_net(x, z_dim, n_z_per_x):
    bn = zs.BayesianNet()
    h = tf.layers.dense(tf.cast(x, tf.float32), 500, activation=tf.nn.relu)
    h = tf.layers.dense(h, 500, activation=tf.nn.relu)
    z_mean = tf.layers.dense(h, z_dim)
    z_logstd = tf.layers.dense(h, z_dim)
    bn.normal("z", z_mean, logstd=z_logstd, group_ndims=1, n_samples=n_z_per_x)
    return bn
```

```
def main():
    # Load MNIST
    data_path = os.path.join(conf.data_dir, "mnist.pkl.gz")
    x_train, t_train, x_valid, t_valid, x_test, t_test = \
```



```
dataset.load_mnist_realval(data_path)
x_train = np.vstack([x_train, x_valid])
x_test = np.random.binomial(1, x_test, size=x_test.shape)
x_dim = x_train.shape[1]

# Define model parameters
z_dim = 40

# Build the computation graph
n_particles = tf.placeholder(tf.int32, shape=[], name="n_particles")
x_input = tf.placeholder(tf.float32, shape=[None, x_dim], name="x")
x = tf.cast(tf.less(tf.random_uniform(tf.shape(x_input)), x_input),
            tf.int32)
n = tf.placeholder(tf.int32, shape=[], name="n")

model = build_gen(x_dim, z_dim, n, n_particles)
variational = build_q_net(x, z_dim, n_particles)

lower_bound = zs.variational.elbo(
    model, {"x": x}, variational=variational, axis=0)
cost = tf.reduce_mean(lower_bound.sgvgb())
lower_bound = tf.reduce_mean(lower_bound)

# # Importance sampling estimates of marginal log likelihood
is_log_likelihood = tf.reduce_mean(
    zs.is_loglikelihood(model, {"x": x}, proposal=variational, axis=0))

optimizer = tf.train.AdamOptimizer(learning_rate=0.001)
infer_op = optimizer.minimize(cost)

# Random generation
x_gen = tf.reshape(model.observe()["x_mean"], [-1, 28, 28, 1])

# Define training/evaluation parameters
epochs = 20
batch_size = 128
iters = x_train.shape[0] // batch_size
save_freq = 10
test_freq = 10
test_batch_size = 400
test_iters = x_test.shape[0] // test_batch_size
result_path = "results/vae"

# Run the inference
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
```

```
for epoch in range(1, epochs + 1):
    time_epoch = -time.time()
    np.random.shuffle(x_train)
    lbs = []
    for t in range(iters):
        x_batch = x_train[t * batch_size:(t + 1) * batch_size]
        _, lb = sess.run([infer_op, lower_bound],
                        feed_dict={x_input: x_batch,
                                   n_particles: 1,
                                   n: batch_size})
        lbs.append(lb)
    time_epoch += time.time()
    print("Epoch {} ({:.1f}s): Lower bound = {}".format(
        epoch, time_epoch, np.mean(lbs)))

    if epoch % test_freq == 0:
        time_test = -time.time()
        test_lbs, test_lls = [], []
        for t in range(test_iters):
            test_x_batch = x_test[t * test_batch_size:
                                   (t + 1) * test_batch_size]
            test_lb = sess.run(lower_bound,
                              feed_dict={x: test_x_batch,
                                           n_particles: 1,
                                           n: test_batch_size})
            test_ll = sess.run(is_log_likelihood,
                              feed_dict={x: test_x_batch,
                                           n_particles: 1000,
                                           n: test_batch_size})
            test_lbs.append(test_lb)
            test_lls.append(test_ll)
        time_test += time.time()
        print(">>> TEST ({:.1f}s)".format(time_test))
        print(">> Test lower bound = {}".format(np.mean(test_lbs)))
        print(">> Test log likelihood (IS) = {}".format(
            np.mean(test_lls)))

    if epoch % save_freq == 0:
        images = sess.run(x_gen, feed_dict={n: 100, n_particles: 1})
        name = os.path.join(result_path,
                             "vae.epoch.{}.png".format(epoch))
        save_image_collections(images, name)

if __name__ == "__main__":
    main()
```